

LAB 2

TA 黃聖祐

Deadline: 2025/10/12 (Sun) 23:59

Demo: 2025/10/13 (Mon)

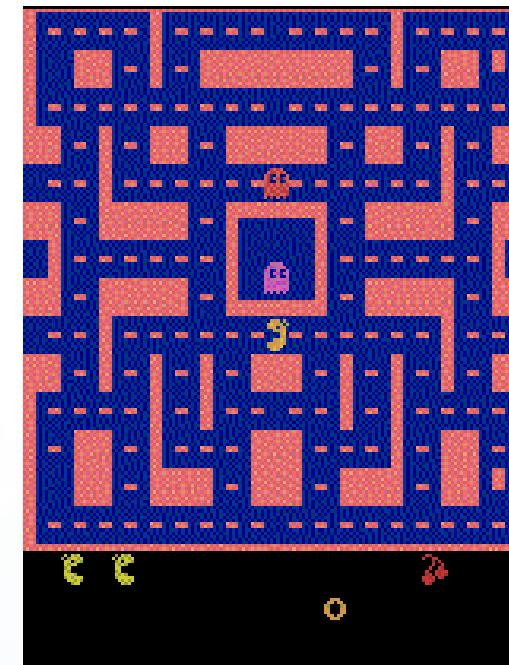
In this lab,

**Must use sample code,
otherwise no credit.**

MsPacman-v5

- Introduction:
 - Your goal is to collect all of the pellets on the screen while avoiding the ghosts.
- Observation space:
 - The whole image
- Action space:

Num	Action
0	NOOP
1	UP
2	RIGHT
3	LEFT
4	DOWN
5	UPRIGHT
6	UPLEFT
7	DOWNRIGHT
8	DOWNLEFT



https://www.gymlibrary.dev/environments/atari/ms_pacman/

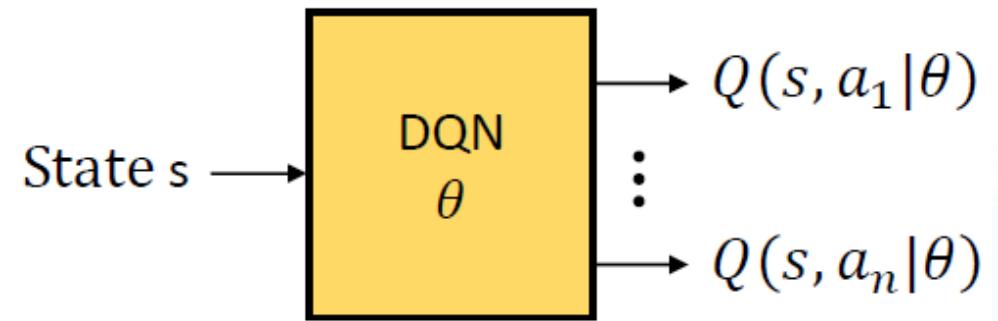
Deep Q-Network (DQN)

Target Q:

$$Y_t^Q = r_{t+1} + \gamma \max_{a'} Q(s_{t+1}, a' | \theta)$$

Loss function:

$$L_Q(s_t, a_t | \theta) = (Y_t^Q - Q(s_t, a_t | \theta))^2$$



Deep Q-Network (DQN)

Algorithm 1 – Deep Q-learning with experience replay:

Initialize replay memory D to capacity N

Initialize action-value function Q with random weights θ

Initialize target action-value function \hat{Q} with weights $\theta^- = \theta$

For episode = 1, M **do**

 Initialize sequence $s_1 = \{x_1\}$ and preprocessed sequence $\phi_1 = \phi(s_1)$

For $t = 1, T$ **do**

 With probability ϵ select a random action a_t

 otherwise select $a_t = \text{argmax}_a Q(\phi(s_t), a; \theta)$

 Execute action a_t in emulator and observe reward r_t and image x_{t+1}

 Set $s_{t+1} = s_t, a_t, x_{t+1}$ and preprocess $\phi_{t+1} = \phi(s_{t+1})$

 Store transition $(\phi_t, a_t, r_t, \phi_{t+1})$ in D

 Sample random minibatch of transitions $(\phi_j, a_j, r_j, \phi_{j+1})$ from D

 Set $y_j = \begin{cases} r_j & \text{if episode terminates at step } j+1 \\ r_j + \gamma \max_{a'} \hat{Q}(\phi_{j+1}, a'; \theta^-) & \text{otherwise} \end{cases}$

 Perform a gradient descent step on $(y_j - Q(\phi_j, a_j; \theta))^2$ with respect to the network parameters θ

 Every C steps reset $\hat{Q} = Q$

End For

End For

Behavior and target network

ϵ -greedy based on behavior network

Experience replay

Update behavior and target network

TODO

- Solve MsPacman-v5 using DQN.
- Find the #TODO comments and hints, remove the raise NotImplementedError.
- Screenshot your Tensorboard training curve and testing results and put it on the report.
- Inherit from the “DQNBaseAgent” and override the “decide_agent_actions” and “update_behavior_network” functions.
- Need to do (please record these in the report)
 - Add Double DQN in your implementation.
 - Add Dueling DQN in your implementation.
 - Add parallelized rollout in your implementation.
 - Solve Enduro-v5 using DQN.

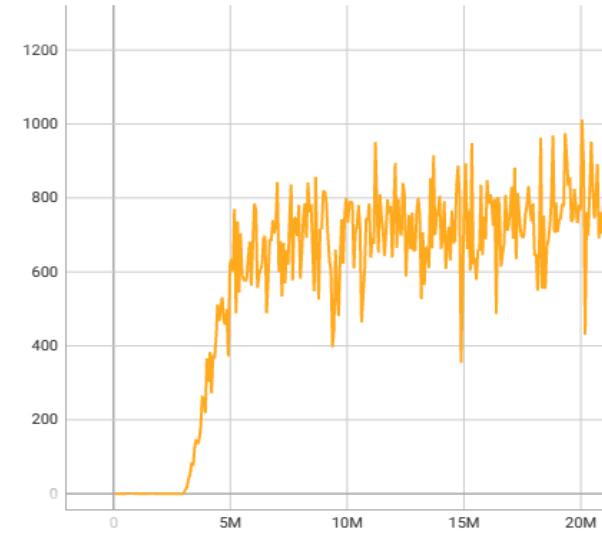
TODO

- Screenshot of Tensorboard training curve and testing results and put it on the report.

Train/Episode Reward



Evaluate/Episode Reward



Training curve:

Testing results (5 games):

Episode: 1	Length: 13301	Total reward: 1040.00
Episode: 2	Length: 13307	Total reward: 1066.00
Episode: 3	Length: 9969	Total reward: 705.00
Episode: 4	Length: 13286	Total reward: 1059.00
Episode: 5	Length: 16630	Total reward: 1318.00
average score: 1037.6		

TODO

- Inherit from the “`DQNBaseAgent`” and override the “`decide_agent_actions`” and “`update_behavior_network`” functions.

base_agent.py

```
class DQNBaseAgent(ABC):
    ...
    @abstractmethod
    def decide_agent_actions(...):
        ...
    def update(...):
        ...
    @abstractmethod
    def update_behavior_network(...):
        ...
    def update_target_network(...):
        ...
    def train(...):
        ...
    def evaluate(...):
        ...
```

dqn_agent_atari.py

```
class AtariDQNAgent(DQNBaseAgent):
    ...
    def decide_agent_actions(...):
        ...
    def update_behavior_network(...):
        ...
```

Double DQN

- DQN suffers from over-estimation.
- Behavior and Target network.
- Reduce the over-estimation problem.



$$Y_t^Q = r_{t+1} + \gamma \max_a Q(S_{t+1}, a | \theta^-)$$



$$Y_t^{DoubleQ} = r_{t+1} + \gamma Q\left(S_{t+1}, \underset{a}{\operatorname{argmax}} Q(S_{t+1}, a | \theta) | \theta^-\right)$$

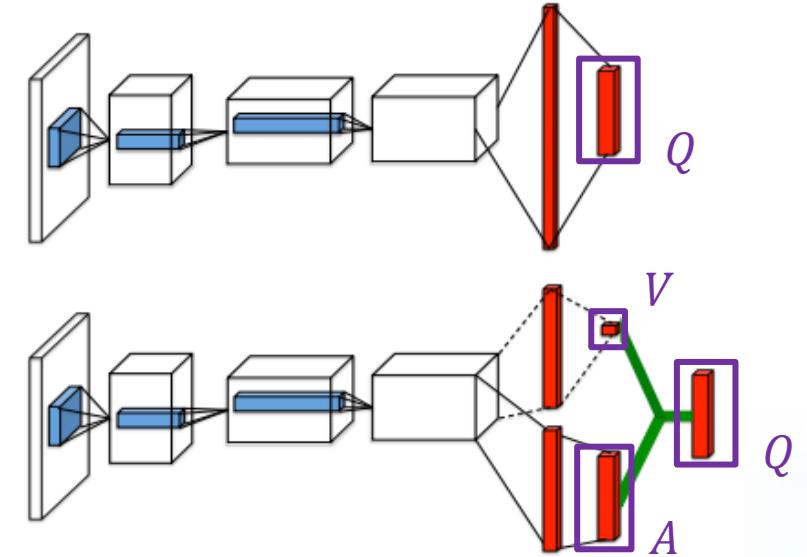
Dueling Networks

- $A(s_t, a_t) = Q(s_t, a_t) - V(s_t)$
→ $Q(s_t, a_t) = A(s_t, a_t) + V(s_t)$

- Constrain the value of A:

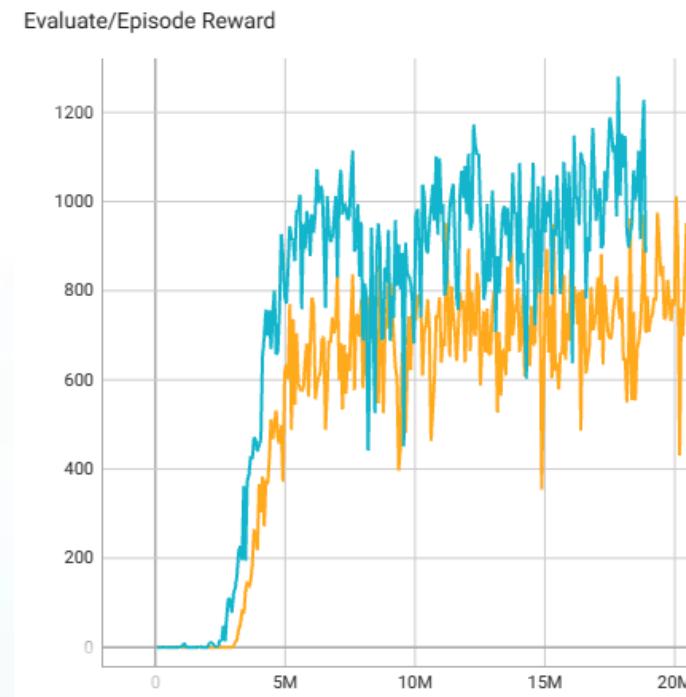
$$Q(s_t, a_t) = V(s) + (A(s_t, a_t) - \frac{1}{|A|} \sum_{a'_t} A(s_t, a'_t))$$

- Due to the relatively small numerical range of the value A, it is more sensitive to model updates, making it easier for the model to consider the **relative changes** in relation to other actions.



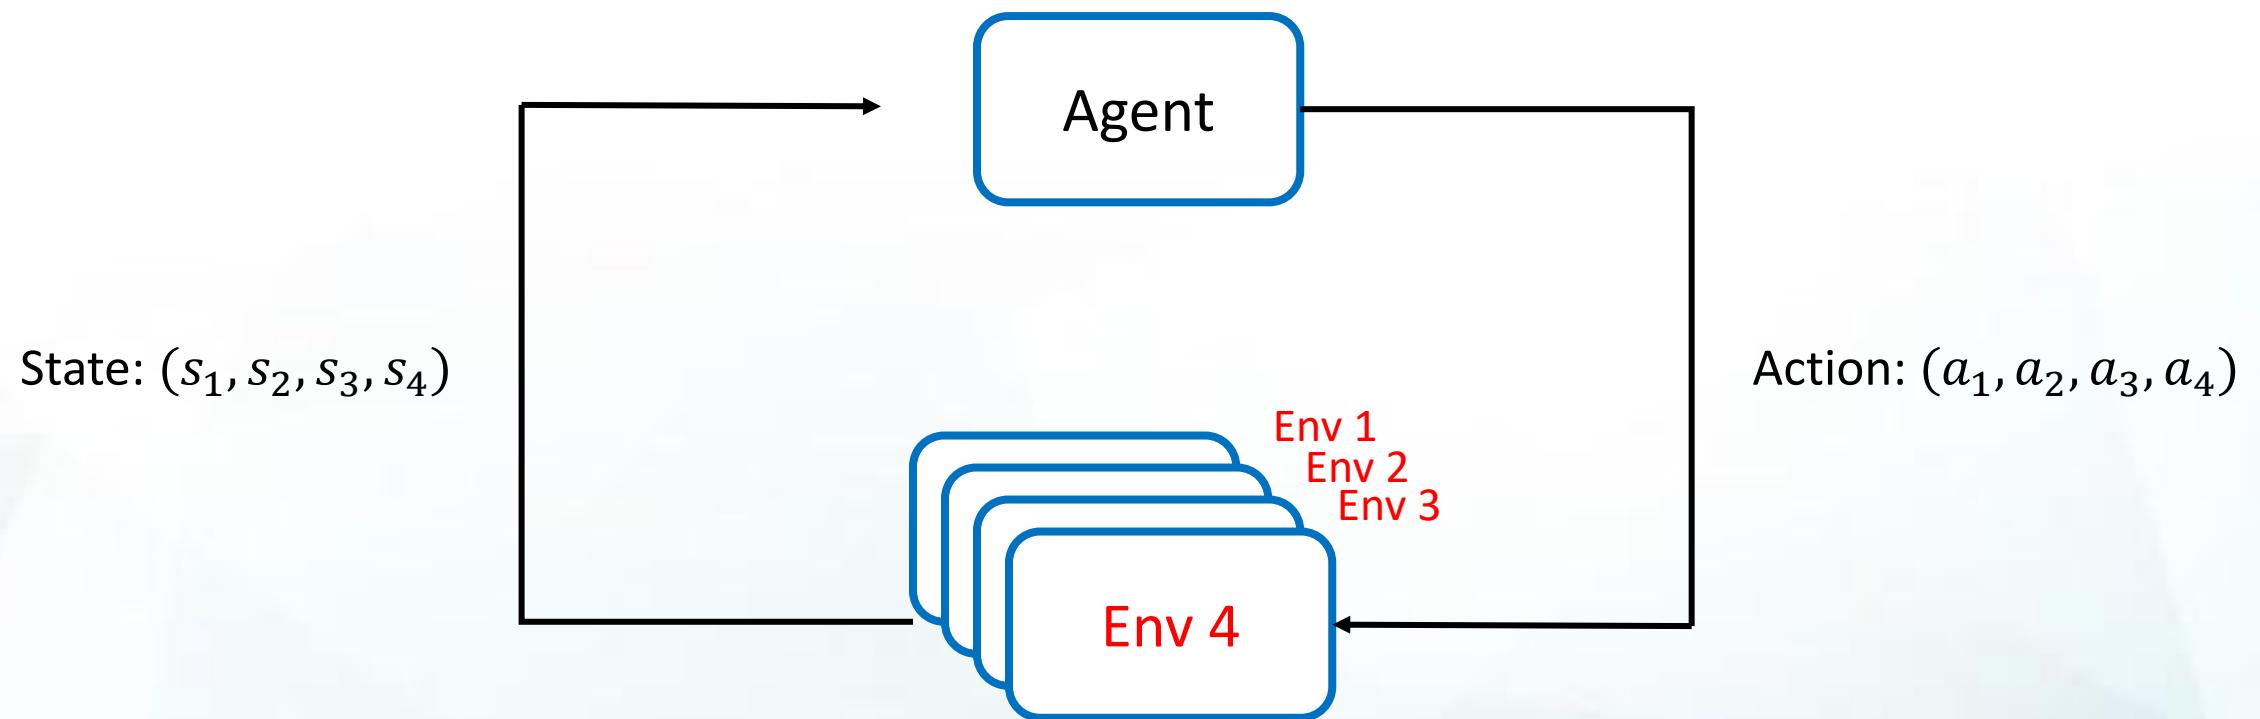
TODO

- Screenshot of Tensorboard training curve and testing results and put it on the report.



Parallelized Rollout

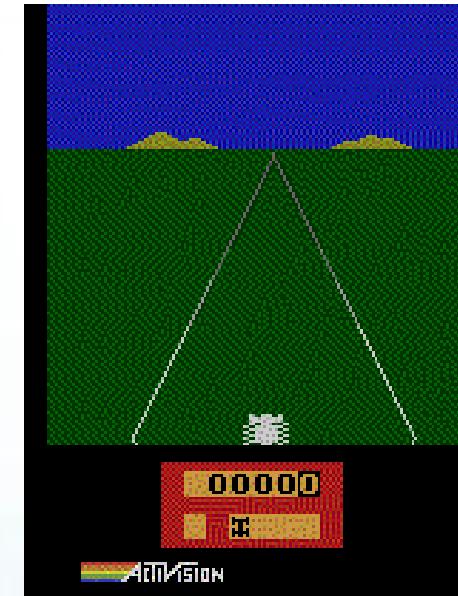
- Parallelly running multiple environments for data collection can help accelerate the training of RL.



Enduro-v5

- Introduction:
 - You are a racer in the National Enduro, a long-distance endurance race. You must overtake a certain amount of cars each day to stay on the race. The first day you need to pass 200 cars, and 300 for each following day. The game ends if you do not meet your overtake quota for the day.
- Observation space:
 - The whole image.
- Action space:

Num	Action
0	NOOP
1	UP
2	RIGHT
3	LEFT
4	DOWN
5	UPRIGHT
6	UPLEFT
7	DOWNRIGHT
8	DOWNLEFT

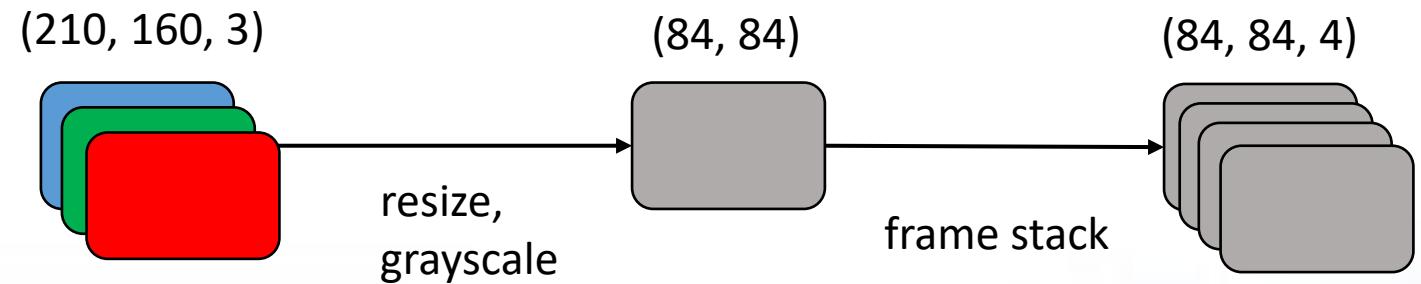


<https://www.gymlibrary.dev/environments/atari/enduro/>

Hint

- You can add any trick you want. E.g.:

- Frame stack
- Clip reward
- Grayscale observation
-



- You can use OpenAI Gym Wrapper.
 - <https://www.gymlibrary.dev/api/wrappers/>

Scoring Criteria

- Your Score = report (30%) + demo performance (50%) + demo questions (20%)
- Report contains two parts:
 - Experimental Results (13%)
 - Screenshot of Tensorboard training curve and testing results on DQN.
 - Experimental Results and Discussion of bonus parts (17%).
 - Screenshot of Tensorboard training curve and testing results on Enduro-v5 (7%).
 - Screenshot of Tensorboard training curve and testing results on DDQN, and discuss the difference between DQN and DDQN (3%).
 - Screenshot of Tensorboard training curve and testing results on Dueling DQN, and discuss the difference between DQN and Dueling DQN (3%).
 - Screenshot of Tensorboard training curve and testing results on DQN with parallelized rollout, and discuss the difference between DQN and DQN with parallelized rollout (4%).

Scoring Criteria

- Screenshot of Tensorboard training curve and testing results and put it on the report.



Scoring Criteria - Demo Performance

- Test your best model for **one game**.
- You have to show the video while testing. You can use env.render() or save video function to achieve this.
- You can use a fixed random seed to reproduce your best game score.



Reward (Game score)	Points (50%)
0 ~ 500	0
500 ~ 1000	10
1000 ~ 1500	20
1500 ~ 2000	30
2000 ~ 2500	40
2500 ~ 3000	45
> 3000	50

Recommended Package Version

- gym 0.26.2
- numpy 1.25.2
- pytorch 2.0.1
- tensorboard 2.14.0
- opencv-python 4.8.0.76
- moviepy 1.0.3

Reminders

- Your network architecture and hyper-parameters **can** differ from the defaults.
- Ensure the **shape** of tensors all the time especially when calculating the **loss**.
- **with no_grad()** : scope is the same as **xxx.detach()**

References

1. Mnih, Volodymyr et al. “Playing Atari with Deep Reinforcement Learning.” ArXiv abs/1312.5602 (2013).
2. Mnih, Volodymyr et al. “Human-level control through deep reinforcement learning.” Nature 518 (2015):529-533.
3. Van Hasselt, Hado, Arthur Guez, and David Silver. “Deep Reinforcement Learning with DoubleQ-Learning.” AAAI. 2016.
4. OpenAI. “OpenAI Gym Documentation.” Retrieved from Getting Started with Gym: <https://gym.openai.com/docs/> .
5. PyTorch. “Reinforcement Learning (DQN) Tutorial.” Retrieved from PyTorch Tutorials: https://pytorch.org/tutorials/intermediate/reinforcement_q_learning.html .