

LAB 4

TA 李謙蓉

Deadline: 2025/11/2 (Sun) 23:59


Demo: 2025/11/3 (Mon) 18:00

In this lab,

**Must use sample code,
otherwise no credit.**

CarRacing-v2

- Introduction:
 - The easiest control task to learn from pixels - a top-down racing environment.
 - The generated track is random every episode.
- Observation space:
 - The whole image
- Action space:
 - Steering (-1 is full left, +1 is full right)
 - Gas (0~1)
 - Braking (0~1)



https://gymnasium.farama.org/environments/box2d/car_racing/

Twin Delayed DDPG (TD3)

- TD3: Add 3 tricks in DDPG
 1. Clipped Double Q-Learning for Actor-Critic
 2. Delayed Policy Updates
 3. Target Policy Smoothing Regularization

Twin Delayed DDPG (TD3)

Algorithm – TD3 algorithm:

Algorithm 1 TD3

```
Initialize critic networks  $Q_{\theta_1}, Q_{\theta_2}$ , and actor network  $\pi_\phi$   
with random parameters  $\theta_1, \theta_2, \phi$   
Initialize target networks  $\theta'_1 \leftarrow \theta_1, \theta'_2 \leftarrow \theta_2, \phi' \leftarrow \phi$   
Initialize replay buffer  $\mathcal{B}$   
for  $t = 1$  to  $T$  do  
    Select action with exploration noise  $a \sim \pi_\phi(s) + \epsilon$ ,  
     $\epsilon \sim \mathcal{N}(0, \sigma)$  and observe reward  $r$  and new state  $s'$   
    Store transition tuple  $(s, a, r, s')$  in  $\mathcal{B}$   
    Sample mini-batch of  $N$  transitions  $(s, a, r, s')$  from  $\mathcal{B}$   
     $\tilde{a} \leftarrow \pi_{\phi'}(s') + \epsilon, \quad \epsilon \sim \text{clip}(\mathcal{N}(0, \tilde{\sigma}), -c, c)$   
     $y \leftarrow r + \gamma \min_{i=1,2} Q_{\theta'_i}(s', \tilde{a})$   
    Update critics  $\theta_i \leftarrow \operatorname{argmin}_{\theta_i} N^{-1} \sum (y - Q_{\theta_i}(s, a))^2$   
    if  $t \bmod d$  then  
        Update  $\phi$  by the deterministic policy gradient:  
         $\nabla_\phi J(\phi) = N^{-1} \sum \nabla_a Q_{\theta_1}(s, a)|_{a=\pi_\phi(s)} \nabla_\phi \pi_\phi(s)$   
        Update target networks:  
         $\theta'_i \leftarrow \tau \theta_i + (1 - \tau) \theta'_i$   
         $\phi' \leftarrow \tau \phi + (1 - \tau) \phi'$   
    end if  
end for
```

1. Clipped Double Q-Learning for Actor-Critic

2. Delayed Policy Updates

3. Target Policy Smoothing Regularization

TODO

- Solve CarRacing-v2 using TD3.
- Ablation study:
 1. Screenshot of Tensorboard training curve and compare the performance of using twin Q-networks and single Q-networks in TD3, and explain (5%).
 2. Screenshot of Tensorboard training curve and compare the impact of enabling and disabling target policy smoothing in TD3, and explain (5%).
 3. Screenshot of Tensorboard training curve and compare the impact of delayed update steps and compare the results, and explain (5%).
 4. Screenshot of Tensorboard training curve and compare the effects of adding different levels of action noise (exploration noise) in TD3, and explain (5%).
 5. Screenshot of Tensorboard training curve and compare your reward function with the original one and explain why your reward function works better (10%).

3 tricks

Exploration
noise


Reward
function

TODO

- Find the #TODO comments and hints, remove the raise NotImplementedError.
- Inherit from the “`TD3BaseAgent`” and override the “`decide_agent_actions`” and “`update_behavior_network`” functions.
- You can try your reward function and network architecture.
- Screenshot of Tensorboard training curve and testing results and put it on the report.

TODO

- Screenshot of Tensorboard training curve and testing results and put it on the report.



Training curve:

Episode: 1	Length: 999	Total reward: 874.44
Episode: 2	Length: 999	Total reward: 883.05
Episode: 3	Length: 999	Total reward: 797.44
Episode: 4	Length: 999	Total reward: 679.18
Episode: 5	Length: 999	Total reward: 866.78
Episode: 6	Length: 999	Total reward: 888.97
Episode: 7	Length: 751	Total reward: 924.80
Episode: 8	Length: 999	Total reward: 883.33
Episode: 9	Length: 999	Total reward: 614.81
Episode: 10	Length: 999	Total reward: 878.34
average score: 829.1142945389436		

Testing results (10 games):

TODO

- Screenshot of Tensorboard training curve and testing results and put it on the report.



Scoring Criteria

Your Score (100%) = report (30%) + demo performance (50%) + demo questions (20%)

- Report contains two parts:
 - Experimental Results (5%)
 - (1) Screenshot of Tensorboard training curve and testing results on TD3.
 - Experimental Results and Discussion (Impact of Twin Q-Networks, Target Policy Smoothing, Delayed Policy Update Mechanism, Action Noise Injection) (25%)
 - (1) Screenshot of Tensorboard training curve and compare the performance of using twin Q-networks and single Q-networks in TD3, and explain (5%).
 - (2) Screenshot of Tensorboard training curve and compare the impact of enabling and disabling target policy smoothing in TD3, and explain (5%).
 - (3) Screenshot of Tensorboard training curve and compare the impact of delayed update steps and compare the results, and explain (5%).
 - (4) Screenshot of Tensorboard training curve and compare the effects of adding different levels of action noise (exploration noise) in TD3, and explain (5%).
 - (5) Screenshot of Tensorboard training curve and compare your reward function with the original one and explain why your reward function works better (5%).

Scoring Criteria

- Screenshot of Tensorboard training curve and testing results and put it on the report.



Scoring Criteria - Demo Performance

- Test your best model for **five** race tracks. Seeds of five tracks will be given on **demo day (3 hours before)**.
- You have to show the video while testing. You can use `env.render()` or save video function to achieve this.



Five race tracks avg.


Reward	Points (50%)
0~100	0
100~199	10
200~299	20
300~399	25
400~499	30
500~599	35
600~699	40
700~799	45
800~	50

Scoring Criteria - Demo Performance


- Test your best model for **five** race tracks. **Seeds of five tracks will be given on demo day (3 hours before).**
- You can decide the race track by giving seed in env.reset().
For example:

```
obs, info = self.env.reset(seed=?)
```

`self.env.reset(seed=2)`



`self.env.reset(seed=5)`



Tensorboard Remote Server

- ssh -p [your port] -L 6006:localhost:6006 pp037@140.113.215.196
- tensorboard --logdir log/dqn
- Open your browser locally and input 127.0.0.1:6006

Recommended Package Version

- gym 0.26.2
- numpy 1.25.2
- pytorch 2.0.1
- tensorboard 2.14.0
- opencv-python 4.8.0.76
- moviepy 1.0.3

Reminders

- Your network architecture and hyper-parameters **can** differ from the defaults.
- Ensure the **shape** of tensors all the time especially when calculating the **loss**.
- **with no_grad()** : scope is the same as **xxx.detach()**
- Be aware of the **indentation** of hints.

References

1. Lillicrap, Timothy P. et al. "Continuous control with deep reinforcement learning." CoRR abs/1509.02971 (2015).
2. Silver, David et al. "Deterministic Policy Gradient Algorithms." ICML (2014).
3. OpenAI. "OpenAI Gym Documentation." Retrieved from Getting Started with Gym: <https://gym.openai.com/docs/>.
4. PyTorch. "Reinforcement Learning (DQN) Tutorial." Retrieved from PyTorch Tutorials: https://pytorch.org/tutorials/intermediate/reinforcement_q_learning.html.
5. Dankwa, Stephen, and Wenfeng Zheng. "Twin-delayed ddpg: A deep reinforcement learning technique to model a continuous movement of an intelligent robot agent." Proceedings of the 3rd international conference on vision, image and signal processing. 2019.
6. My results: https://youtu.be/FAqATf_k5fI?si=p_bwyjt4RDchLmJ5